OpenBench — Presenter Pack

Script · numbers cheat sheet · Q&A bank · risk card

Team Inferno · Tech Zephyr Case Finale, IIT Bhubaneswar · 31 October 2026. Contents: (1) speaker script with timings, (2) one-page numbers cheat sheet, (3) Q&A bank, (4) risk and trigger card. Every number comes from OpenBench_Financial_Model.xlsx.

1 · Speaker script (10-minute slot, target 9:20)

Slide 1 · **Mahi** · 14s · ends at 0:14

The Idle Instrument Paradox

Good morning. We are Team Inferno, and we were asked whether OpenBench can scale trust, not just bookings. Our short answer: yes, but only if it stops selling access to instruments and starts selling delivered results.

Slide 2 · **Mahi** · 53s · ends at 1:07

Sell delivered results, not listings, and take ₹1.2 Cr in two tranches

Here is our answer up front. OpenBench should sell delivered results, not listings. Concretely, three moves. One: a managed tier we call Assured, for the pilot labs, the deep-tech belt and the startup hub, where the customer gets a committed slot, sealed custody and an on-time guarantee. That takes on-time from 78% today to about 92% by Month 12. Two: we grow in a deliberate order and reach 17 host labs by Month 12, entering the price-sensitive manufacturing corridor last. Three: in our base case OpenBench breaks even in Month 15 and uses only about ₹38.6 lakh of the ₹1.2 crore. On the investor, we renegotiate rather than reject: same money, paid in two tranches, with two targets lowered and a stricter quality bar added. The rest of the deck shows why.

Slide 3 · **Mahi** · 50s · ends at 1:57

Customers leave because supply is listed, not productised: three structural causes

The pilot shows four symptoms: 14% of orders cancelled by the lab, 8% late, only 42% of customers coming back, and lab managers who feel underpaid. We trace these to three structural causes. First, supply is listed but not productised: only 38 of 86 instruments are genuinely bookable, and nothing commits a staffed slot behind a listing. Second, labs earn too little per hassle, so when teaching clashes with an external job, the external job loses. Third, there is no trust layer; 31% of customers cite unpredictable turnaround and 27% worry about confidentiality. And this is not an OpenBench quirk. India's own free portal, I-STEM, has executed only about half of its booking requests. So a better listing will not fix this. A guaranteed outcome will.

Slide 4 · **Mahi** · 42s · ends at 2:39

At an 18% fee the pilot loses ₹2.95L monthly and needs 5.5× its volume

Now the money. On an average ₹8,500 booking, OpenBench keeps an 18% fee, ₹1,530, and spends ₹600 supporting it, leaving ₹930. Seventy bookings a month gives ₹65,100 of contribution against ₹3.6 lakh of fixed cost, so the pilot loses about ₹2.95 lakh every month. To break even on today's economics it would need about 387 bookings a month, 5.5 times today's volume. The tempting move is to chase volume. But with 14% of orders cancelled and only 42% of customers returning, more volume mostly means more broken promises. So our sequence is reliability first, then scale. Minakshi will now show which model does that.

Slide 5 · **Minakshi Singha** · 66s · ends at 3:45

A segmented hybrid beats every pure model on reliability and investor fit

We scored all three models, plus a hybrid, on six weighted criteria. Reliability carries the most weight, 25 points, because it is what the investor's 90% target and the customer's repeat decision both depend on. The pure marketplace is cheapest to onboard but can only reach 82% on-time. Managed fulfilment reaches 94% but costs ₹1.2 lakh per lab and is expensive for price-sensitive buyers. Lab SaaS fails on adoption: only 4 of 10 lab managers would pay. Our hybrid scores 82. One backbone serves all tiers; only Assured runs before the Month-7 gate. The customer promise is simple: booked means delivered. The customer pays per booking; OpenBench keeps 28% on Assured jobs, or 8% plus a lab subscription on Connect. OpenBench runs everything around the instrument. And labs join because they earn more, not less: a pilot lab's monthly income rises about 5.6%, because committed slots fill and it keeps half of every priority premium, while we take the admin away and pay within seven days.

Slide 6 · **Minakshi Singha** · 65s · ends at 4:50

Enter where trust pays most first; take the price-sensitive corridor last

Sequencing trades value against volume and ease. We start by converting the four pilot labs to Assured in Month 1. Cluster A, the deep-tech belt, comes next: it has the highest booking value, about ₹48,240 per lab per month on Assured, and its key concern, IP, is exactly what our custody layer solves. Cluster C, the startup hub, follows; it serves similar customers, so we reuse the same playbook. That gives us 13 labs in three cities by Month 8, above the investor's 12. The manufacturing corridor, B, has the most volume but the thinnest margin and is price-sensitive, so we enter it last, only after we pass a Month-7 reliability gate, and on a cheaper tier. We do not treat Exhibit 2 as guaranteed: we assume 85% of projected volume, anchored on the pilot running at 87.5% of baseline, and a four-month ramp for every new lab. Most customers come free through incubators and I-STEM leads; the rest cost ₹1,500 each.

Slide 7 · **Minakshi Singha** · 66s · ends at 5:56

Base case breaks even in Month 15; downside keeps ₹41.4L cash but misses M18

Here is the 18-month model; every figure links back to one assumptions tab in our Excel file. The teal line is the base case: we lose up to ₹3.6 lakh a month while onboarding, and cross break-even in Month 15, three months ahead of the investor's deadline. We deliver 2,316 paid bookings in the first year and use about ₹38.6 lakh. The red line is our downside: lower capture, lower prices, a two-month delay and cost overruns. It does not break even by Month 18, but it never runs short of cash: we still hold about ₹41.4 lakh at the lowest point. The box on the right shows what matters most: booking value, capture rate and fixed cost. And we are transparent: break-even depends on two operating levers, a priority-SLA premium and cost learning in fulfilment. Our 50% priority premium matches published express rates. Switch both off and Month-18 EBITDA is −₹1.23L. In the downside, our Month-9 response narrows the loss to ₹2.32 lakh.

Slide 8 · **Minakshi Singha** · 54s · ends at 6:50

Accept the money, renegotiate two targets, and give the investor a tighter quality bar

We tested every investor condition against our own model rather than promising what we cannot deliver. We beat the host target: 17 labs by Month 12. We beat 90% on-time and break-even by Month 18. Two asks we would miss honestly: 2,316 bookings against 2,400, and a 58% repeat rate against 60%, because repeat lags reliability by a quarter. So we propose 2,200 bookings and 55% repeat, while keeping 60% for Assured customers. Published evidence points to 51 to 55%. We also propose paying the ₹1.2 crore in two halves: ₹60 lakh now, ₹60 lakh at Month 7 if we have ten labs live and 90% on-time. Even our downside only needs ₹37.6 lakh by then. The investor gives up nothing: less money at risk, a stricter quality covenant, a board observer and monthly reporting.

Slide 9 · **Swapnil Behera** · 54s · ends at 7:44

A six-step trust layer turns a listed instrument into a guaranteed result

This is what customers pay 28% for. Every job moves through six controlled steps. Labs commit staffed slot blocks, and only those appear as bookable, which closes the 86-versus-38 gap. Samples travel under a QR chain of custody with photo logs. Calibration is checked before every test, and expired instruments delist themselves. The client owns the data under a three-party NDA. Payment sits in escrow and reaches the lab within seven days. And disputes go to mediation, then a neutral re-test. When something breaks, the rules are automatic: a cancelled job reroutes to a backup lab and the host loses 10% of its payout; a late result refunds 20% of our fee. Compare that with a listing portal, where nobody is accountable and, on India's own portal, only about half of requests are carried out.

Slide 10 · **Swapnil Behera** · 48s · ends at 8:32

Four checkpoints and five KPIs show the board early whether it works

Here is how the board will know, early, whether this is working. By Month 3, 6 labs are live and the trust layer is running. By Month 6, 10 labs and on-time at 94%, which sets up the Month-7 tranche gate. By Month 12, 17 labs in four cities and about 2,316 bookings. By Month 18, all 20 labs and a positive monthly EBITDA. We track five KPIs against today's baseline. The two that matter most are on-time, from 78% to at least 90%, and host cancellations, from 14% to at most 5%. We also track the bookable share of listed instruments, 90-day repeat, and contribution per booking, which rises from ₹930 today to about ₹1,354 in Month 18.

Slide 11 · **Swapnil Behera** · 58s · ends at 9:30

Every major risk has a numeric trigger and a pre-agreed response

Finally, risks. We list the six that could break this plan, each with a number that tells us it is happening and a response agreed in advance. The biggest is that labs refuse to commit slots; it happened to Science Exchange when a Stanford school declined to list its facilities. If fewer than ten labs are live at Month 7, we hold the second tranche. If on-time falls below 88% at Month 6, we stop opening cities. If demand runs below 70% of projections by Month 9, we skip the cheapest tier and cut a city's cost. Disintermediation is managed by making the platform worth staying on: escrow, the guarantee and consolidated invoicing. To close: OpenBench wins by selling a delivered result. We ask the board to approve the Assured-first plan and the two-tranche counter-offer. Booked should mean delivered. Thank you; we welcome your questions.

Total: 1424 words ≈ 9.5 minutes at 150 words per minute. Split: Mahi slides 1–4, Minakshi Singha 5–8, Swapnil Behera 9–11 (change freely). Rehearse twice with a stopwatch; cut slide 5's second half first if you run long.

2 · Numbers cheat sheet (memorise these)

M15
break-even

2,316
bookings by M12

92%
on-time at M12

₹38.6L
cash used, base

Number	Value	Where it comes from
Pilot contribution per booking	₹930	18% × ₹8,500 − ₹600 (Ex.1)
Pilot monthly loss	₹2.95L	70 × ₹930 − ₹3.60L fixed
Pilot break-even volume	387 bookings/month (5.5× today)	₹3.60L ÷ ₹930
Bookable share of listed instruments	44% (38 of 86)	Ex.1
I-STEM execution rate	≈50% (44,084 of 88,887)	Office of the PSA, 31 Dec 2025
Contribution per booking by model (baseline)	Marketplace ₹930 · Managed ₹1,030 · SaaS ₹280 + ₹12k/host	Ex.3
Contribution per host at mature volume	A ₹48,240 (Assured) · C ₹23,140 (Assured) · B ₹15,600 (Connect) / ₹17,100 (Verified)	Unit Economics tab
Connect vs Verified crossover in B	26.7 bookings/host (89% of Ex.2)	₹12,000 ÷ (₹570 − ₹120)
Pilot lab income, today vs Assured	₹121,975 → ₹128,775 (+5.6%)	17.5 × 82% vs 20 × 75.75% of ₹8,500 (72% + half the Priority premium)
Scoring matrix	Marketplace 56 · Managed 77 · SaaS 48 · Hybrid 82	Slide 4 weights
Break-even month (base)	M15	Model: Scenarios / Monthly tabs
Paid bookings M1–M12 (base)	2,316	Model: Scenarios / Monthly tabs
Hosts live at M12	17	Model: Scenarios / Monthly tabs
On-time at M12	92.3%	Model: Scenarios / Monthly tabs
90-day repeat at M12	57.8%	Model: Scenarios / Monthly tabs
EBITDA at M18 (base)	+₹0.15L	Model: Scenarios / Monthly tabs
Cash used M1–M18 (base)	₹38.6L	Model: Scenarios / Monthly tabs
Cash used M1–M18 (downside)	₹78.6L	Model: Scenarios / Monthly tabs
EBITDA at M18 (downside)	−₹2.65L	Model: Scenarios / Monthly tabs
Pilot break-even bookings/month	387	Model: Scenarios / Monthly tabs
Hosts by milestone	M3 6 · M6 10 · M12 17 · M18 20	Monthly_Base
Cash low point, downside	₹41.4L	Monthly_Downside
Cash used by M7, downside	₹37.6L (< ₹60L tranche 1)	Monthly_Downside
Upside break-even	M10	Monthly_Upside
Both levers off (no Priority SLA, no cost learning)	M18 EBITDA −₹1.23L	Sensitivity tab
Priority premium	+50% price, 15% uptake, OpenBench keeps half	A22–A24; Auriga 50% express surcharge [S9]
Customer-acquisition cost (CAC)	₹1,500 per paid-acquired org; 80% won free; ≈₹30,127/month at M12; ₹4.6L over 18 months	A31–A33; B2B partnership CAC ₹50k–1.5L [S10]
New customer orgs needed at M12	100 a month	0.405 orgs per booking (170 / 420, Ex.1)

Number	Value	Where it comes from
Lab onboarding funnel	267 contacted → 80 meetings → 20 signed → 16 live (≈21 contacts/month)	Funnel tab; I-STEM 3,812 institutes [S1], 72 AICs [S11], SATHI [S12]
Slow onboarding (12 new labs)	No break-even by M18 (M18 −₹0.32L); lowest cash ₹81.0L	Funnel tab
Downside + M9 response	M18 −₹2.32L; lowest cash ₹47.7L; still no break-even	Monthly_DownResponse
Repeat link A27 at 0.5× / 1.0× / 1.5×	M12 repeat 49.9% / 57.8% / 65.7%	Repeat check tab; evidence 0.55–0.83× [S13]

3 · Q&A bank (answer in ≤30 seconds, always with a number)

Question	Answer
Why not a pure marketplace? It's cheapest to onboard.	It is cheapest (₹45k per lab, 2.4-month payback), but Exhibit 3 caps it at 82% on-time, below the investor's 90%. Cheap supply that cancels 14% of orders is what keeps repeat at 42%.
Why would a university accept your penalties?	Penalties apply only to slots the lab itself committed. In return the lab earns more, not less: a pilot lab goes from ₹121,975 to ₹128,775 a month (+5.6%), with no admin and payment in 7 days. DST's SRIMAN guidelines also reward sharing through incentives and ratings.
Isn't 28% too expensive for Cluster B?	Yes, so we don't charge it there. B gets Connect: customers pay 8%, the lab pays ₹12k a month. Connect beats the marketplace whenever a lab does fewer than 26.7 bookings a month, which is likely in a price-sensitive cluster.
What if only 8 labs sign by Month 7?	Then the second ₹60L tranche is not released, we stop at the cities already opened, and we never add the ₹0.5L-per-city fixed cost for cities we don't enter. The investor's exposure stays capped at ₹60L.
Why should the investor accept a lower bookings target?	Our base delivers 2,316, so 2,400 would be a promise we can't model. 2,200 leaves a 5% cushion. In exchange we accept a stricter quality bar: on-time ≥90% for three straight months and cancellations ≤5%.
How do you hit 90% on-time without owning instruments?	By controlling everything around the instrument: committed staffed slots, backup-lab rerouting and payout penalties. About 79% of Month-12 bookings run on the Assured tier, which Exhibit 3 puts at 94%. Our model gives 92.3% at M12.
Why isn't I-STEM enough? It's free.	It's a listing: it executed about half of its 88,887 requests, and users settle payment and reports directly with the lab. We sell the outcome, and we use I-STEM as our map for finding host labs.
How do you stop customers going direct to the lab?	Leakage is worst on high-value repeat work, so we make staying worth it: escrow, the SLA guarantee and rebates, backup labs, NDA paperwork done for them, one GST invoice across labs, and repeat-customer pricing.
What happens to break-even if fixed costs double?	At ₹11L a month, M18 EBITDA falls to about −₹5.3L; we would not open City 4. Our ₹0.5L-per-city cost is lean because Assured's ₹1,350-per-booking variable cost already pays for fulfilment labour.
Who is liable if a sample is lost or IP leaks?	The chain-of-custody log shows exactly where it happened. The 3-party agreement makes the at-fault party liable, capped at the booking value and insured; the customer owns all data and nobody keeps copies.
How will you win customers in a new city?	Through incubators and startup hubs (Cluster C shares Cluster A's customers) and through industry associations in the corridor. The pilot already has 170 paying organisations. We now model it: 80% of new organisations come free through incubators, SATHI centres, I-STEM leads and referrals, and the rest cost ₹1,500 each (A31–A33), about ₹0.30L a month at M12. The city manager in the ₹0.5L city cost owns this.
Why only partly use SaaS when it has the lowest variable cost?	Because 6 of 10 lab managers refuse a subscription and its contribution at the baseline is the lowest, ₹17,600 per lab. We use it only where its subscription floor protects margin: price-sensitive Cluster B.
Why take ₹1.2 Cr if the base case uses ₹38.6L?	Because the downside uses ₹78.6L. Tranching keeps the unused half with the investor until we prove reliability, and anything left funds the next city only after break-even.
How confident are you in 85% capture?	It is anchored on the pilot, which runs at 87.5% of the Exhibit-3 baseline. Our downside uses 65%. If capture is 70%, M18 EBITDA is about −₹0.58L, which is why capture is a Month-9 trigger.

Question	Answer
Is break-even real, or does it depend on your levers?	It depends on two levers and we say so: the priority-SLA premium and an 11% cut in fulfilment cost. With both off, M18 EBITDA is −₹1.23L. The 50% premium matches a NABL lab's published express surcharge [S9]; both are tracked monthly.
Will universities cap how much external work they take?	Some do: Stanford's medical school kept non-academic users to about 15% of its core facilities' business. That is why labs commit slot blocks they choose themselves, sized inside their own policy.
What is your customer-acquisition cost?	₹1,500 per organisation we pay to win, through association events and webinars (A31); 80% come free through incubators, SATHI centres, I-STEM leads and pilot referrals (A33). That is about ₹30,127 a month at M12 and ₹4.6L over 18 months, earned back in about one booking. If free channels drop to 60%, M18 EBITDA falls to −₹0.28L.
How do you find 16 new labs?	A funnel (appendix A7): 267 labs contacted, 80 meetings, 20 signed, 16 live, about 21 contacts a month. Leads come from I-STEM's 3,812 listed institutes, 72 Atal Incubation Centres and 5 SATHI centres. If only 12 sign, we miss M18 break-even (−₹0.32L) but cash never drops below ₹81.0L.
Why will customers pay a priority premium?	31% of pilot customers cite turnaround as their main concern (Ex.1), and an Indian NABL lab already charges a 50% surcharge for express testing [S9]. We assume only 15% uptake and share half the premium with the lab. Some IIT rate cards list no urgent fee, so we track uptake monthly.
What happens if the downside hits?	Downside: no break-even by M18, −₹2.65L a month, cash floor ₹41.4L. At M9 the demand trigger fires: we skip Cluster B Verified labs and cut City-4 cost. That narrows the loss to −₹2.32L and lifts the floor to ₹47.7L (appendix A9). It does not rescue M18, which is why tranche 2 sits behind the M7 gate.
Why 55% repeat instead of 60%?	The case gives one data point (78% on-time, 42% repeat). A supply-chain study shows 0.55–0.83 points of retention per point of on-time [S13]; at that slope Month-12 repeat is 51–55%. Our model's 1.0× gives 57.8%. 55% is a target we can defend.
What about seasonality in the startup hub?	Cluster C is only 4 of 20 labs. Committed slot blocks are set quarterly, so labs release capacity during exam and teaching peaks rather than cancelling.

4 · Risk and trigger card

Risk	Trigger	Pre-agreed response
Labs refuse to commit slots	<10 hosts live at M7	Hold tranche 2; pitch SRIMAN credit and higher lab income
Reliability slips	On-time <88% at M6	Freeze new cities until fixed
Demand below plan	Bookings/host <70% of Exhibit 2 at M9	Skip B Verified hosts; cut City-4 cost (₹0.5L/month)
Disintermediation	Repeat customers' share of GMV falls 2 months running	Escrow, guarantee, consolidated invoicing, repeat pricing
Price pushback in B	B host on-time <85% or churn >10%	Move host to Connect, or exit
Cash	Cash <₹40L before M12	Stop onboarding; renegotiate tranche timing

KPIs to quote: on-time 78% → ≥90% from M7 · host cancellations 14% → ≤5% by M6 · bookable share 44% → ≥80% · repeat 42% → ≥55% at M12 · contribution per booking ₹930 → ≥₹1,250 at M18.